

# Data Wrangling

```
library(tidyverse)
```

## The Tidyverse

The Tidyverse is a collection of R packages designed for data science that share a common philosophy, grammar, and data structures. It makes working with data in R more consistent and readable.

### Core Tidyverse Packages

Package	Purpose
ggplot2	Data visualisation
dplyr	Data manipulation (filtering, summarising, transforming)
tidyr	Tidying data (reshaping, pivoting)
purrr	Functional programming / advanced iteration
readr	Reading rectangular data files (CSV, TSV, etc.)
tibble	Modern version of <code>data.frame</code> with better printing and subsetting
forcats	Working with factors
stringr	Working with character strings

## dplyr Basics

dplyr aims to provide a function for each basic verb of data manipulation. These verbs can be organised into three categories based on the component of the dataset that they work with:

### *Rows:*

- `filter()` chooses rows based on column values.
- `slice()` chooses rows based on location.
- `arrange()` changes the order of the rows.

### *Columns:*

- `select()` changes whether or not a column is included.
- `rename()` changes the name of columns.
- `mutate()` changes the values of columns and creates new columns.
- `relocate()` changes the order of the columns.

### *Groups of rows:*

- `summarise()` collapses a group into a single row.
- The pipe

All of the dplyr functions take a data frame (or tibble) as the first argument.

At the most basic level, you can only alter a tidy data frame in five useful ways: you can reorder the rows (`arrange()`), pick observations and variables of interest (`filter()` and `select()`), add new variables that are functions of existing variables (`mutate()`), or collapse many values to a summary (`summarise()`).

### *Others*

- Pipeline operator `%>%`
- `group_by()`
- Family of `join` operations

## Example

```
name <- c("A", "B", "C", "D", "E", "F")
rcomputing <- c(45,40,90,94,80,65)
eda <- c(63,66,75,83,80,59)
sl <- c(59,56,91,92,86,67)
marks <- data.frame(name,rcomputing,eda,sl)
head(marks, 3)
```

```
##   name rcomputing eda sl
## 1    A          45 63 59
## 2    B          40 66 56
## 3    C          90 75 91
```

## select() function: Select columns

The `select()` function is used to choose specific columns from a data frame (or tibble). This is useful when you want to focus on certain variables or drop the ones you don't need.

```
# Select specific columns by name
select(marks, name, eda)
```

```
# Equivalent using the pipe
marks %>% select(name, eda)
```

```
# Selecting by position
select(marks, 1)           # select col 1
select(marks, 1:3)         # select cols 1 through 3
select(marks, c(1, 3))     # select cols 1 and 3
select(marks, -1)          # select all cols except column 1
```

```
marks %>% select(name:eda)  # select all cols between name & eda
marks %>% select(!(name:eda)) # select all cols except those between name & eda
```

- Column names act as positions in `select()`.
  - For example, `select(marks, name)` is equivalent to `select(marks, 1)` if `name` is the first column.
- Because of this, variables from the surrounding environment are not automatically used inside `select()`.

```
rprogramming <- 5
select(marks, rprogramming) # looks for a col called 'rprogramming'
select(marks, 5)             # error if column 5 doesn't exist
```

## Contextual variables in helpers

This restriction only applies to bare names (e.g., `name`, `eda`, `name:eda`). When using selection helpers (like `starts_with()`, `ends_with()`, etc.), external variables can be used.

```
specify <- "ing"
select(marks, ends_with(specify)) # selects all columns ending with "ing"
```

You can also select columns based on specific criteria with:

- `starts_with()` - select cols that start with a character string
- `ends_with()` - select cols that end with a character string
- `contains()` - select cols that contain a character string
- `matches()` - select cols that match a regular expression
- `one_of()` - select cols names that are from a group of names

```
select(marks, starts_with("n"))
marks %>% select(ends_with("ing"))
```

## filter() function: Filter rows

- The `filter()` function is used to extract a subset of rows from a data frame (or tibble).
- You provide conditions, and only the rows where those conditions evaluate to `TRUE` are returned.

```
# Select rows where rcomputing > 50
filter(marks, rcomputing > 50)

# Multiple conditions (rows where both are TRUE)
marks %>% filter(rcomputing > 50, eda > 75)
```

This is roughly equivalent to base R subsetting:

```
marks[marks$rcomputing > 50 & marks$eda > 75, ]
```

## Equality and logical expressions

```
filter(marks, eda == 80) # rows where eda is exactly 80
filter(marks, rcomputing > 50 & eda > 60) # multiple logical conditions
```

## Using between()

The `between()` helper checks if a value falls within a closed interval (inclusive).

```
filter(marks, between(rcomputing, 50, 80))
```

This is the same as:

```
filter(marks, rcomputing >= 50 & rcomputing <= 80)
```

## arrange() function: Arrange rows

The `arrange()` function reorders the rows of a data frame (or tibble) according to one or more variables.

```
# Arrange rows by eda (ascending, the default)
arrange(marks, eda)

# Equivalent with pipes
marks %>% arrange(eda, sl)
```

- By default, sorting is ascending.
- When multiple columns are given, subsequent columns are used to break ties.
  - In the example above, rows with the same `eda` value are ordered by `sl`.

Use `desc()` to sort in descending order:

```
marks %>% arrange(desc(eda))
```

## slice() function: Choose rows by position

The `slice()` family of functions selects rows based on their integer positions, rather than conditions on values (as with `filter()`). This makes it useful for sampling, keeping top/bottom rows, or indexing directly.

```
marks %>% slice(3:5) # select rows at positions 3 to 5
```

First or last rows:

```
marks %>% slice_head(n = 2) # first 2 rows
marks %>% slice_tail(n = 2) # last 2 rows
```

Random rows:

```
marks %>% slice_sample(n = 3) # randomly select 3 rows
marks %>% slice_sample(prop = 0.7) # randomly select 70% of rows
```

- Use `replace = TRUE` for bootstrap sampling.
- Add the `weight_by` argument to perform weighted sampling.

Highest or lowest values:

- Use `slice_min()` and `slice_max()` to select rows with the smallest or largest values of a variable.

- Remove missing values first if necessary.

```
marks %>%
  filter(!is.na(eda)) %>%
  slice_max(eda, n = 3) # top 3 rows by eda
```

## rename() function: Change column names

The `rename()` function changes the names of existing columns.

The syntax is:

```
rename(.data, new_name = old_name)
```

```
# Rename column 'sl' to 'supervised'
rename(marks, supervised = sl)
```

The new name goes on the left, the old name on the right.

Use `rename_with()` when you want to apply a function to rename multiple columns at once (e.g. converting all names to lowercase).

```
marks %>% rename_with(tolower) # all names to lowercase
marks %>% rename_with(~ gsub(" ", "_", .x)) # replace spaces with underscores
```

You can rename columns selected with `starts_with()`, `ends_with()`, `matches()`, etc.

```
marks %>% rename_with(toupper, starts_with("r"))
# all columns starting with "r" converted to UPPERCASE
```

If the new names are stored in a variable, use `setNames()` inside `rename_with()`.

```
new_names <- c("student_name", "exploratory", "supervised")
marks %>% rename_with(~ new_names)
```

## mutate() function: Add or transform columns

The `mutate()` function creates new variables or transforms existing ones in a data frame (or tibble).

- Unlike `select()`, which works with column names/positions, `mutate()` works with column vectors (the actual values).
- Often used for feature engineering: standardizing, creating flags, applying mathematical/logical transformations.
- Order aware: later mutations can use columns created earlier in the same call.
- Supports logical conditions, boolean operators, and helper arguments like `.keep`.

1. Create new variables

```
mutate(marks, st.eda = (eda - mean(eda)) / sd(eda)) # standardize eda
```

2. Conditional variables

```
marks <- mutate(marks, pass.rcomp = ifelse(rcomputing < 50, "fail", "pass"))
```

3. Chain multiple mutations

```
marks %>%
  mutate(pass.rcomp = ifelse(rcomputing < 50, "fail", "pass"),
         log_eda = log(eda))
```

4. Keep only new variables

Use `.keep = "none"` to return just the new column(s):

```
mutate(marks, pass.rcomp = ifelse(rcomputing < 50, "fail", "pass"),
       .keep = "none")
```

5. Use existing mutations immediately

```
marks %>%
  mutate(total = rcomputing + eda,
         avg   = total / 2)  # 'avg' uses 'total' created above
```

## 6. Logical and boolean transformations

```
marks %>%
  mutate(high_score = eda > 75 & rcomputing > 60)
```

## 7. Programmatic renaming with across()

Combine `mutate()` with `across()` for applying transformations to multiple columns at once:

```
marks %>%
  mutate(across(c(eda, rcomputing), log, .names = "log_{.col}"))
```

## relocate() function: Change column order

The `relocate()` function reorders columns in a data frame (or tibble).

```
relocate(.data, ..., .before = NULL, .after = NULL)
```

- `.data`: the data frame or tibble.
- `...`: columns to move (by name, range, or helper functions).
- `.before`: move the selected columns before a given column.
- `.after`: move the selected columns after a given column.

```
marks %>% relocate(rcomputing:eda, .before = name)
```

Move to front or back

```
marks %>% relocate(eda)  # move 'eda' to the front
marks %>% relocate(eda, .after = last_col())  # move 'eda' to the end
```

Reorder multiple columns

```
marks %>% relocate(c(name, sl), .before = rcomputing)
```

Use helper functions

```
marks %>% relocate(starts_with("r"), .after = name)  # move all 'r*' cols after 'name'
```

## summarise() function: Summarise values

The `summarise()` function collapses a data frame into a single row or grouped summaries by computing summary statistics such as mean, sum, min, max, count, etc.

```
marks %>% summarise(eda_mean = mean(eda, na.rm = TRUE))
```

```
##   eda_mean
## 1      71
```

## Summarise all numeric columns

`summarise_all()` applies a summary function to every column.

```
marks %>%
  select_if(is.numeric) %>%          # keep numeric columns
  summarise_all(~ sum(., na.rm = TRUE)) # sum each column, ignoring NA
```

- `~` defines a formula.
- `.` represents the current column's data.

## Using across() (modern approach)

`summarise_all()` is superseded. The recommended way is to use `across()` inside `summarise()`.

```
marks %>%
  summarise(across(where(is.numeric), ~ sum(., na.rm = TRUE)))
```

Multiple summaries at once

```
marks %>%
  summarise(across(where(is.numeric),
    list(mean = ~ mean(., na.rm = TRUE),
         sd    = ~ sd(., na.rm = TRUE))))
```

## Grouped summaries

Combine with `group_by()` to summarise by category:

```
marks %>%
  group_by(pass.rcomp) %>%
  summarise(eda_mean = mean(eda, na.rm = TRUE),
            count = n())
```

## The pipeline operator %>%

The pipeline operator (`%>%`) from `magrittr` (and used heavily in `dplyr`) allows you to chain multiple operations in a clear, readable sequence.

- It passes the output of one function as the first argument of the next function.
- This avoids nested function calls like: `third(second(first(x)))`
- Instead, we write: `x %>% first() %>% second() %>% third()`

```
marks %>%
  mutate(year = 2020) %>%
  filter(eda > 60)
```

Multi-step workflows

```
marks %>%
  filter(eda > 60) %>%
  mutate(pass = ifelse(rcomputing > 50, "pass", "fail")) %>%
  arrange(desc(eda)) %>%
  summarise(avg_eda = mean(eda, na.rm = TRUE))
```

## `group_by()` function: Group data by a variable

The `group_by()` function is used to create groups within a data frame, so that subsequent operations (e.g. `summarise()`, `mutate()`, `filter()`) are applied within each group rather than across the entire dataset.

- `group_by()` does not change the data immediately — it just defines the grouping structure.
- Use `ungroup()` to remove the grouping after operations are complete.

```
marks %>%
  group_by(pass.rcomp) %>%
  summarise(eda_mean = mean(eda, na.rm = TRUE)) # separate mean for each pass.rcomp value
```

Inspect groups:

```
marks %>% distinct(pass.rcomp) # see the unique group values
```

Grouped mutation:

```
marks %>%
  group_by(pass.rcomp) %>%
  mutate(max_eda = max(eda, na.rm = TRUE)) # col with max eda in each group
```

Ungroup after operations:

```
marks %>%
  group_by(pass.rcomp) %>%
```

```
summarise(eda_mean = mean(eda, na.rm = TRUE)) %>%
ungroup() # ensures further operations are applied to the entire dataset
```

Multiple grouping variables:

```
marks %>%
  group_by(pass.rcomp, year) %>%
  summarise(avg_eda = mean(eda, na.rm = TRUE))
```

## count() function: Count occurrences

The `count()` function counts the number of occurrences of each unique value in a column, or a combination of columns.

By default, it returns:

1. A column with the unique values from the specified variable(s)
2. A column `n` showing the frequency of each value

```
marks %>% count(pass.rcomp)
# Counts how many rows belong to each value of pass.rcomp
```

Count multiple columns:

```
marks %>% count(pass.rcomp, year)
# Counts occurrences for combinations of values across multiple columns
```

Sort counts:

```
marks %>% count(pass.rcomp, sort = TRUE)
# Orders results by n in descending order (most frequent first)
```

Weighted counts:

```
marks %>% count(pass.rcomp, wt = eda)
# Instead of counting rows, sums the eda values within each group
```

Equivalent `group_by()` + `summarise()`:

```
marks %>%
  group_by(pass.rcomp) %>%
  summarise(n = n())
# Shows that count() is a shortcut for grouped summarisation.
```

## Family of join operations

The `dplyr` package provides a set of join functions to merge two data frames (or tibbles) based on one or more key columns.

Syntax:

```
joined_data <- join_type(df1, df2, by = "key_column")
```

- `by` specifies the column(s) used to match rows between the two tibbles.
- If the column names are identical in both tables, `by` can be omitted.
- If the columns have different names, use a named vector:

```
by = c(df1_col = "df2_col")
```

```
df1 <- tibble(id = c(1, 2, 3), name = c("Alice", "Bob", "Charlie"))
df2 <- tibble(id = c(2, 3, 4), score = c(90, 85, 88))
```

```
inner_join(df1, df2, by = "id")
left_join(df1, df2, by = "id")
right_join(df1, df2, by = "id")
full_join(df1, df2, by = "id")
semi_join(df1, df2, by = "id")
anti_join(df1, df2, by = "id")
```

Joining on multiple columns:

```
full_join(df1, df2, by = c("col1", "col2", "col3"))
```

- Combines rows where all specified columns match.
- Non-matching rows are kept with NA in columns that don't exist in the other table.

Joining columns with different names:

```
full_join(df1, df2, by = c("col1" = "colX", "col2" = "colY"))  
# Maps col1 in df1 to colX in df2, and col2 in df1 to colY in df2
```

Note:

- Always check column names before joining.
- Use `distinct()` if necessary to remove duplicates before joining.
- Use `select()` to include only the relevant key columns if the datasets are wide

## Ranking Functions in dplyr

Ranking functions assign relative positions to numeric values, with different strategies for handling ties. These are useful for ranking performance scores, sales, or any numeric variable.

### 1. min\_rank()

- Assigns ranks to values.
- Tied values receive the same rank.
- The smallest possible rank is assigned to tied values, and subsequent ranks are skipped.

```
x <- c(10, 20, 20, 30)  
min_rank(x)  
# [1] 1 2 2 4
```

### 2. dense\_rank()

- Similar to `min_rank()`.
- Tied values receive the same rank, but no ranks are skipped.

```
dense_rank(x)  
# [1] 1 2 2 3
```

### 3. row\_number()

- Assigns unique ranks to each value, even when tied.

```
row_number(x)  
# [1] 1 2 3 4
```

### 4. percent\_rank()

- Computes the relative percentile rank of each value, ranging from 0 to 1.
- 0 corresponds to the smallest value, 1 to the largest.

```
percent_rank(x)  
# [1] 0.0 0.3333 0.3333 1.0
```

### 5. ntile()

- Divides data into n quantiles (e.g. quartiles, deciles).
- Assigns each value a rank corresponding to its quantile.

```
ntile(x, 2) # Divide into 2 quantiles  
# [1] 1 1 2 2
```



Function	Tie Handling	Output Type	Notes
<code>min_rank()</code>	Smallest rank for ties, skips next	Integer rank	Skips numbers after ties
<code>dense_rank()</code>	Same rank for ties, no skipping	Integer rank	Consecutive ranks
<code>row_number()</code>	Unique rank for each value	Integer rank	Tied values get arbitrary order
<code>percent_rank()</code>	Same rank for ties	Numeric 0–1	Relative percentile
<code>ntile(n)</code>	Groups into n quantiles	Integer rank 1–n	Useful for categories

## NYCFlights13 Dataset

On-time data for all flights that departed NYC (i.e. JFK, LGA or EWR) in 2013.

- year, month, day Date of departure.
- dep\_time, arr\_time: Actual departure and arrival times (format HHMM or HMM), local tz.
- sched\_dep\_time, sched\_arr\_time: Scheduled departure and arrival times (format HHMM or HMM), local tz.
- dep\_delay, arr\_delay: Departure and arrival delays, in minutes. Negative times represent early departures/arrivals.
- carrier: Two letter carrier abbreviation. See airlines to get name.
- flight: Flight number.
- tailnum: Plane tail number. See planes for additional metadata.
- origin, dest: Origin and destination. See airports for additional metadata.
- air\_time: Amount of time spent in the air, in minutes.
- distance: Distance between airports, in miles.
- hour, minute: Time of scheduled departure broken into hour and minutes.
- time\_hour: Scheduled date and hour of the flight as a POSIXct date. Along with origin, can be used to join flights data to weather data.

```
library(nycflights13)
head(flights, 3)
```

```
## # A tibble: 3 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1     1     517             515         2     830             819
## 2  2013     1     1     533             529         4     850             830
## 3  2013     1     1     542             540         2     923             850
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

1. Find all flights that

a) Had an arrival delay of two or more hours

```
filter(flights, arr_delay >= 120)
```

b) Flew to Houston (IAH or HOU)

```
filter(flights, dest == "IAH" | dest == "HOU")
```

c) Were operated by United, American, or Delta

```
filter(flights, carrier %in% c("UA", "AA", "DL"))
# OR
filter(flights, carrier == "UA" | carrier == "AA" | carrier == "DL")
```

d) Departed in summer (July, August, and September)

```
filter(flights, month %in% c(7, 8, 9))
```

e) Arrived more than two hours late, but didn't leave late

```
filter(flights, arr_delay > 120 & dep_delay == 0)
```

f) Were delayed by at least an hour, but made up over 30 minutes in flight

```
filter(flights, dep_delay >= 60 & arr_delay <= 30)
```

g) Departed between midnight and 6am (inclusive)

```
filter(flights, dep_time >= 0 & dep_time <= 600)
```

2. Use between() to simplify the code needed to answer the previous questions.

```
filter(flights, between(dep_time, 0, 600))
filter(flights, between(month, 7, 9))
```

3. How many flights have a missing dep\_time? What other variables are missing? What might these rows represent?

```
nrow(filter(flights, is.na(dep_time)))
```

```
flights %>%  
  summarise_all(~ sum(is.na(.)))
```

4. Sort the flights according to day, month and year

```
arrange(flights, day, month, year)
```

5. Sort the flights using the arrival time in a descending order.

```
arrange(flights, desc(arr_time))
```

6. How could you use arrange() to sort all missing values to the start?

```
arrange(flights, desc(is.na(arr_time)))
```

7. Sort flights to find the most delayed flights. Find the flights that left earliest.

```
arrange(flights, desc(dep_delay)) # most delayed flights at the top  
arrange(flights, dep_delay) # flights that left the earliest (neg delay times)
```

8. Sort flights to find the fastest (highest speed) flights.

```
arrange(flights, air_time/distance)
```

9. Which flights travelled the farthest? Which travelled the shortest?

```
head(arrange(flights, desc(distance)), 3)  
head(arrange(flights, distance), 3)
```

10. Select all columns in the flights dataframe between year and day (inclusive).

```
select(flights, year:day)
```

11. Select all columns except those from year to day (inclusive).

```
select(flights, (!year:day))
```

12. Rename the tail\_num variable in flights dataframe with tailnum.

```
rename(flights, tailnum = tail_num)
```

13. Using the pipeline operator do the following:

a) Select all columns in the flights dataframe between year and day (inclusive).

```
flights %>%  
  select(year:day)
```

b) Select all columns that ends with delay and time.

```
flights %>%  
  select(ends_with("delay"), ends_with("time"))
```

c) Select the distance and air\_time variables.

```
flights %>%  
  select(distance, air_time)
```

d) Create a gain/loss travel time for each flight.

```
flights %>%  
  mutate(gain_loss = ifelse((arr_time - sched_arr_time) > 0, "gain",  
                           ifelse((arr_time - sched_arr_time) < 0, "loss", "no gain/loss")))
```

e) What is the speed of the flight.

```
flights %>%  
  mutate(speed = distance / (air_time / 60)) # convert air_time to hours
```

14. Convert `dep_time` and `sched_dep_time` to minutes since midnight.

- The `dep_time` and `sched_dep_time` columns are convenient to look at (format HHMM), but they are not continuous numeric values, which makes calculations tricky.
- We can convert them to minutes since midnight using integer division (`%/%`) and modulus (`%%`).

```
flights %>%
  mutate(
    dep_time_minutes = (dep_time %/% 100) * 60 + (dep_time %% 100),
    sched_dep_time_minutes = (sched_dep_time %/% 100) * 60 + (sched_dep_time %% 100))
```

- `dep_time %/% 100` → extracts the hour part (integer division by 100)
- `dep_time %% 100` → extracts the minute part (remainder after division by 100)
- Multiply the hour part by 60 and add the minutes to get total minutes since midnight

Example of `%/%` and `%%`

```
123 %/% 10 # 12 - quotient (tens)
123 %% 10  # 3 - remainder (ones)
```

- `%/%` → integer division (quotient)
- `%%` → modulus (remainder)

15. Compare `air_time` with `arr_time - dep_time`. What do you expect to see? What do you see? What do you need to do to fix it?

```
flights %>%
  mutate(actual_air_time = arr_time - dep_time) %>%
  select(actual_air_time, air_time)
```

- `dep_time` and `arr_time` are HHMM numbers, not minutes since midnight.
- Subtracting HHMM numbers does not account for hours and minutes correctly.
- Flights that cross midnight or are delayed make the subtraction even more misleading.

How to fix it:

1. Convert `dep_time` and `arr_time` to minutes since midnight
2. If `arr_time_minutes < dep_time_minutes` (flight crosses midnight), add  $24 \times 60$  to `arr_time_minutes` before subtraction:

```
flights %>%
  mutate(
    dep_time_minutes = (dep_time %/% 100) * 60 + (dep_time %% 100),
    arr_time_minutes = (arr_time %/% 100) * 60 + (arr_time %% 100),
    actual_air_time = ifelse(arr_time_minutes < dep_time_minutes,
                           arr_time_minutes + 1440 - dep_time_minutes,
                           arr_time_minutes - dep_time_minutes)
  ) %>%
  select(actual_air_time, air_time)
```

## Tidying Data: `pivot_longer()` & `pivot_wider()`

The `tidyr` package provides functions to reshape data between long and wide formats.

### 1. `pivot_longer()`

- Converts data from **wide format** (many columns) into **long format** (fewer columns, more rows).
- Useful when you have multiple columns representing the same type of measurement.

```
pivot_longer(
  data,
  cols,           # columns to pivot
  names_to,       # name of the new key column
  values_to)     # name of the new value column
```

```
df <- tibble(id = 1:2, math = c(90, 80), english = c(85, 75))
df

## # A tibble: 2 x 3
##       id  math english
##   <int> <dbl>   <dbl>
## 1     1    90     85
## 2     2    80     75

df_long <- df %>%
  pivot_longer(cols = math:english, names_to = "subject", values_to = "score")
df_long

## # A tibble: 4 x 3
##       id subject score
##   <int> <chr>   <dbl>
## 1     1  math     90
## 2     1 english    85
## 3     2  math     80
## 4     2 english    75
```

## 2. pivot\_wider()

- Converts data from long format into wide format.
- Useful when you want separate columns for different values of a variable.

```
pivot_wider(
  data,
  names_from,    # column whose values become new column names
  values_from)  # column containing values to fill
```

```
df_wide <- df_long %>%
  pivot_wider(
    names_from = subject,
    values_from = score)
df_wide
```

```
## # A tibble: 2 x 3
##       id  math english
##   <int> <dbl>   <dbl>
## 1     1    90     85
## 2     2    80     75
```